

Boosting Algorithms Project (AdaBoost, Gradient Boosting, XGBoost)

This notebook is a mini-project you can upload to GitHub.

Goals:

1. Understand boosting by training and comparing:

- Decision Stump (weak learner baseline)
- AdaBoost (adaptive boosting)
- Gradient Boosting (stage-wise additive model)
- XGBoost (optimized gradient boosting)

2. Produce clean, reproducible outputs:

- Metrics: accuracy, balanced accuracy, F1
- Confusion matrix + classification report
- Decision boundary plots (2D)
- Feature importance (where applicable)

Dataset:

- Wine dataset (from scikit-learn)

```
In [1]: # Important libraries to install
#!pip install -U numpy pandas matplotlib scikit-learn
#!pip install -U xgboost
```

```
In [2]: # Numerical computing
import numpy as np

# Data handling (tables)
import pandas as pd

# Plotting
import matplotlib.pyplot as plt

# scikit-learn datasets
from sklearn.datasets import load_wine

# Splitting data into train/test sets
from sklearn.model_selection import train_test_split

# Metrics for evaluation
from sklearn.metrics import (
```

```

accuracy_score,
balanced_accuracy_score,
f1_score,
classification_report,
confusion_matrix,
ConfusionMatrixDisplay
)

# Tree model (used as a decision stump)
from sklearn.tree import DecisionTreeClassifier

# Boosting models (AdaBoost + Gradient Boosting)
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClass

# For consistent randomness across runs
RANDOM_STATE = 42

```

In [3]:

```

def evaluate_classifier(model, X_train, X_test, y_train, y_test, title
    """
    Train a model, predict, and print useful metrics.

    Parameters
    -----
    model : estimator object
        Any scikit-learn compatible classifier with fit/predict.
    X_train, X_test : array-like
        Feature matrices for training and test.
    y_train, y_test : array-like
        Target labels for training and test.
    title : str
        Header name in printed output.

    Returns
    -----
    fitted_model : estimator object
        The same model after calling .fit(...)
    y_pred_test : np.ndarray
        Predictions on X_test
    """
    # Fit (train) the model on training data
    model.fit(X_train, y_train)

    # Predict labels for test data
    y_pred_test = model.predict(X_test)

    # Compute metrics
    acc = accuracy_score(y_test, y_pred_test)
    bacc = balanced_accuracy_score(y_test, y_pred_test)
    f1 = f1_score(y_test, y_pred_test, average="macro") # macro treat

    # Print the results
    print(f"\n{title} ===")

```

```

print(f"Accuracy : {acc:.4f}")
print(f"Balanced Accuracy : {bacc:.4f}")
print(f"Macro F1 : {f1:.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_test))

return model, y_pred_test

def plot_confusion(y_true, y_pred, title="Confusion Matrix"):
    """
    Plot a confusion matrix using scikit-learn's built-in display help
    """
    cm = confusion_matrix(y_true, y_pred)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)
    disp.plot(values_format="d")
    plt.title(title)
    plt.show()

def plot_decision_boundary_2d(model, X, y, feature_names=("x1", "x2"),
                               title="2D decision regions for a classifier.
    Works ONLY when X has exactly 2 columns.
    """
    # Determine min/max range for each feature axis
    x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
    y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5

    # Create a dense grid of points covering the feature space
    xx, yy = np.meshgrid(
        np.linspace(x_min, x_max, 400),
        np.linspace(y_min, y_max, 400)
    )

    # Stack grid into shape (n_points, 2) for prediction
    grid = np.c_[xx.ravel(), yy.ravel()]

    # Predict labels for every point in the grid
    preds = model.predict(grid).reshape(xx.shape)

    # Plot the decision regions
    plt.figure(figsize=(8, 4))
    plt.contourf(xx, yy, preds, alpha=0.25)

    # Plot training points on top
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolor="k", s=35)

    # Labels and title
    plt.xlabel(feature_names[0])
    plt.ylabel(feature_names[1])

```

```
plt.title(title)
plt.show()
```

```
In [4]: # Helper Functions (Evaluation and Plotting)
def evaluate_classifier(model, X_train, X_test, y_train, y_test, title
    """
        Train a model, predict, and print useful metrics.

    Parameters
    -----
    model : estimator object
        Any scikit-learn compatible classifier with fit/predict.
    X_train, X_test : array-like
        Feature matrices for training and test.
    y_train, y_test : array-like
        Target labels for training and test.
    title : str
        Header name in printed output.

    Returns
    -----
    fitted_model : estimator object
        The same model after calling .fit(...)
    y_pred_test : np.ndarray
        Predictions on X_test
    """
    # Fit (train) the model on training data
    model.fit(X_train, y_train)

    # Predict labels for test data
    y_pred_test = model.predict(X_test)

    # Compute metrics
    acc = accuracy_score(y_test, y_pred_test)
    bacc = balanced_accuracy_score(y_test, y_pred_test)
    f1 = f1_score(y_test, y_pred_test, average="macro") # macro treat

    # Print the results
    print(f"\n{title} ===")
    print(f"Accuracy      : {acc:.4f}")
    print(f"Balanced Accuracy : {bacc:.4f}")
    print(f"Macro F1       : {f1:.4f}")
    print("\nClassification Report:")
    print(classification_report(y_test, y_pred_test))

    return model, y_pred_test

def plot_confusion(y_true, y_pred, title="Confusion Matrix"):
    """
        Plot a confusion matrix using scikit-learn's built-in display help
    """
```

```

cm = confusion_matrix(y_true, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(values_format="d")
plt.title(title)
plt.show()

def plot_decision_boundary_2d(model, X, y, feature_names=("x1", "x2"),
.....  

    Plot 2D decision regions for a classifier.  

    Works ONLY when X has exactly 2 columns.  

.....
# Determine min/max range for each feature axis
x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5

# Create a dense grid of points covering the feature space
xx, yy = np.meshgrid(
    np.linspace(x_min, x_max, 400),
    np.linspace(y_min, y_max, 400)
)

# Stack grid into shape (n_points, 2) for prediction
grid = np.c_[xx.ravel(), yy.ravel()]

# Predict labels for every point in the grid
preds = model.predict(grid).reshape(xx.shape)

# Plot the decision regions
plt.figure(figsize=(8, 4))
plt.contourf(xx, yy, preds, alpha=0.25)

# Plot training points on top
plt.scatter(X[:, 0], X[:, 1], c=y, edgecolor="k", s=35)

# Labels and title
plt.xlabel(feature_names[0])
plt.ylabel(feature_names[1])
plt.title(title)
plt.show()

```

In [5]:

```

# Part A: AdaBoosting vs Decision Stump (2D Visualization)
# Load wine dataset as a pandas-friendly object
wine = load_wine(as_frame=True)

# Create a DataFrame containing all features
X_full = wine.data.copy()

# Target labels (0,1,2 correspond to wine classes)
y_full = wine.target.copy()

```

```

# Feature names (for reference)
feature_names = wine.feature_names
print("Feature count:", X_full.shape[1])
print("First 5 features:", feature_names[:5])

# Select two features to match the textbook-style plot:
# - "alcohol"
# - "od280/od315_of_diluted_wines"
X_2d = X_full[["alcohol", "od280/od315_of_diluted_wines"]].copy()

# Convert to NumPy arrays for plotting and modeling
X_2d_np = X_2d.values
y_np = y_full.values

```

Feature count: 13

First 5 features: ['alcohol', 'malic_acid', 'ash', 'alcalinity_of_ash', 'magnesium']

```

In [6]: # Make it a binary Classification problem (Cleaner Boundaries)
        # For clearer boosting explanation + 2D decision boundary,
        # we restrict the dataset to TWO classes (binary classification).
        # We'll keep classes 0 and 1 only.

        mask_binary = (y_np == 0) | (y_np == 1)
        X_bin = X_2d_np[mask_binary]
        y_bin = y_np[mask_binary]

        print("Binary dataset shape:", X_bin.shape)
        print("Class counts:", np.bincount(y_bin))

```

Binary dataset shape: (130, 2)

Class counts: [59 71]

```

In [7]: # Test/Train Split
        # Split into train and test sets.
        # stratify=y_bin ensures both sets keep similar class proportions.

        X_train, X_test, y_train, y_test = train_test_split(
            X_bin, y_bin,
            test_size=0.30,
            random_state=RANDOM_STATE,
            stratify=y_bin
        )

        print("Train shape:", X_train.shape, "Test shape:", X_test.shape)

```

Train shape: (91, 2) Test shape: (39, 2)

```

In [8]: # Weak Learner Baseline: Decision Stump (Depth=1)
        # Split into train and test sets.
        # stratify=y_bin ensures both sets keep similar class proportions.

        X_train, X_test, y_train, y_test = train_test_split(
            X_bin, y_bin,

```

```

        test_size=0.30,
        random_state=RANDOM_STATE,
        stratify=y_bin
    )

print("Train shape:", X_train.shape, "Test shape:", X_test.shape)

```

Train shape: (91, 2) Test shape: (39, 2)

```

In [9]: # AdaBoosting with Decision Stump (Adaptive Boosting)
# AdaBoost trains stumps sequentially.
# Each new stump focuses more on examples that were misclassified prev
# (via sample weights).

ada = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1, random_state=RANDOM_
    n_estimators=200,           # number of weak learners (stumps)
    learning_rate=0.1,          # shrinkage factor (smaller can be more st
    random_state=RANDOM_STATE
)

ada, pred_ada = evaluate_classifier(
    ada, X_train, X_test, y_train, y_test,
    title="AdaBoost (stumps)"
)

plot_confusion(y_test, pred_ada, title="AdaBoost - Confusion Matrix")

plot_decision_boundary_2d(
    ada, X_train, y_train,
    feature_names=("Alcohol", "OD280/OD315"),
    title="Decision Boundary - AdaBoost (stumps)"
)

```

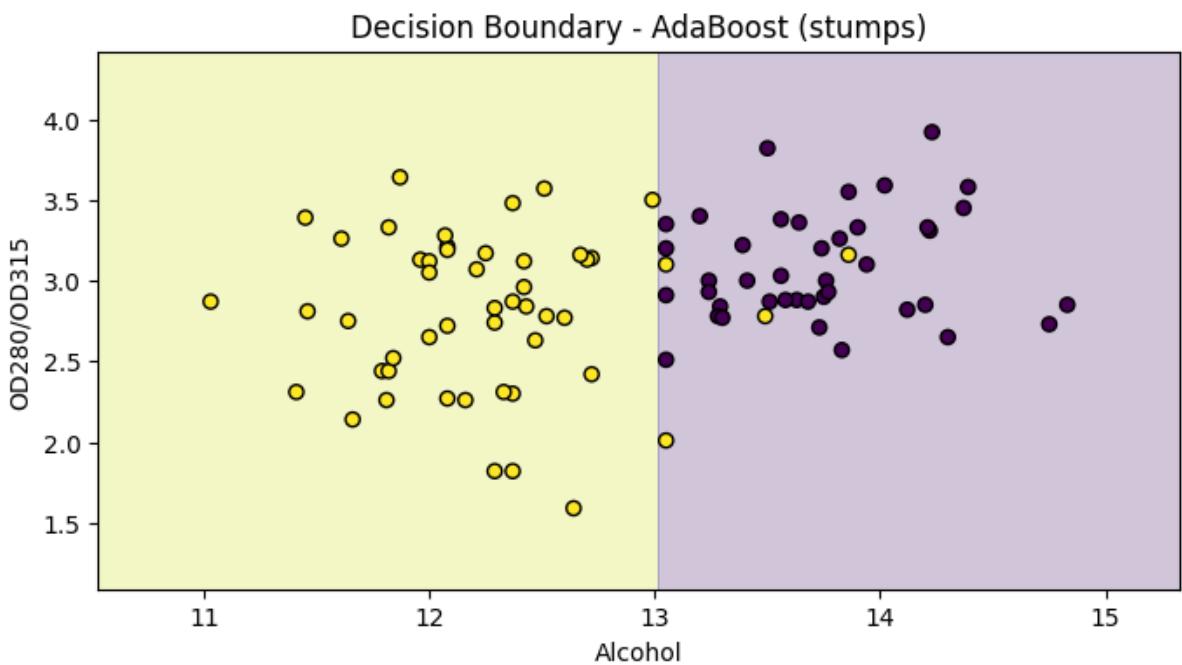
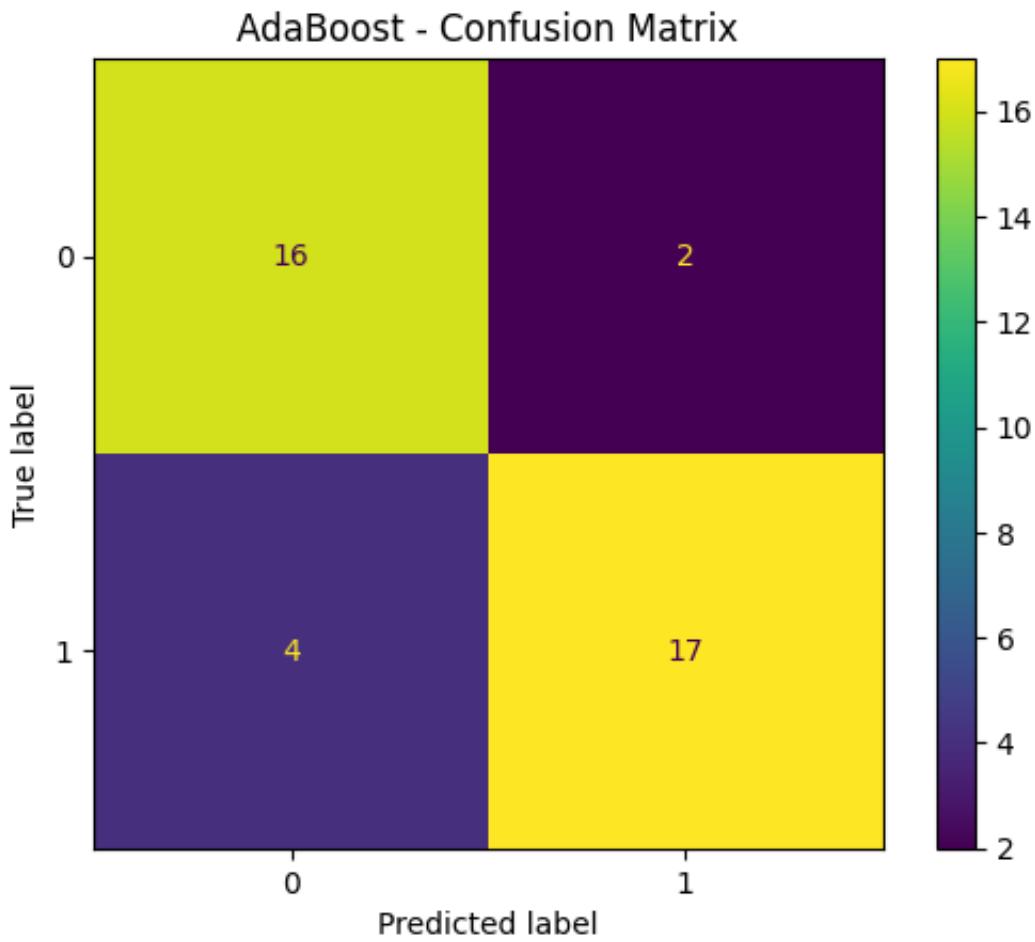
```

==== AdaBoost (stumps) ====
Accuracy : 0.8462
Balanced Accuracy : 0.8492
Macro F1 : 0.8461

```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.89	0.84	18
1	0.89	0.81	0.85	21
accuracy			0.85	39
macro avg	0.85	0.85	0.85	39
weighted avg	0.85	0.85	0.85	39



```
In [10]: # Part B: Gradient Boosting
# GradientBoostingClassifier
# Gradient Boosting builds an additive model in a stage-wise fashion.
# Each new tree tries to reduce the errors made by the current ensemble

gb = GradientBoostingClassifier(
```

```

        n_estimators=200,           # number of boosting stages
        learning_rate=0.1,          # shrink contribution of each tree
        max_depth=3,                # depth of each individual tree
        random_state=RANDOM_STATE
    )

gb, pred_gb = evaluate_classifier(
    gb, X_train, X_test, y_train, y_test,
    title="GradientBoostingClassifier (2D, binary)"
)

plot_confusion(y_test, pred_gb, title="GradientBoosting - Confusion Ma

```

```

plot_decision_boundary_2d(
    gb, X_train, y_train,
    feature_names=("Alcohol", "OD280/OD315"),
    title="Decision Boundary - Gradient Boosting"
)

```

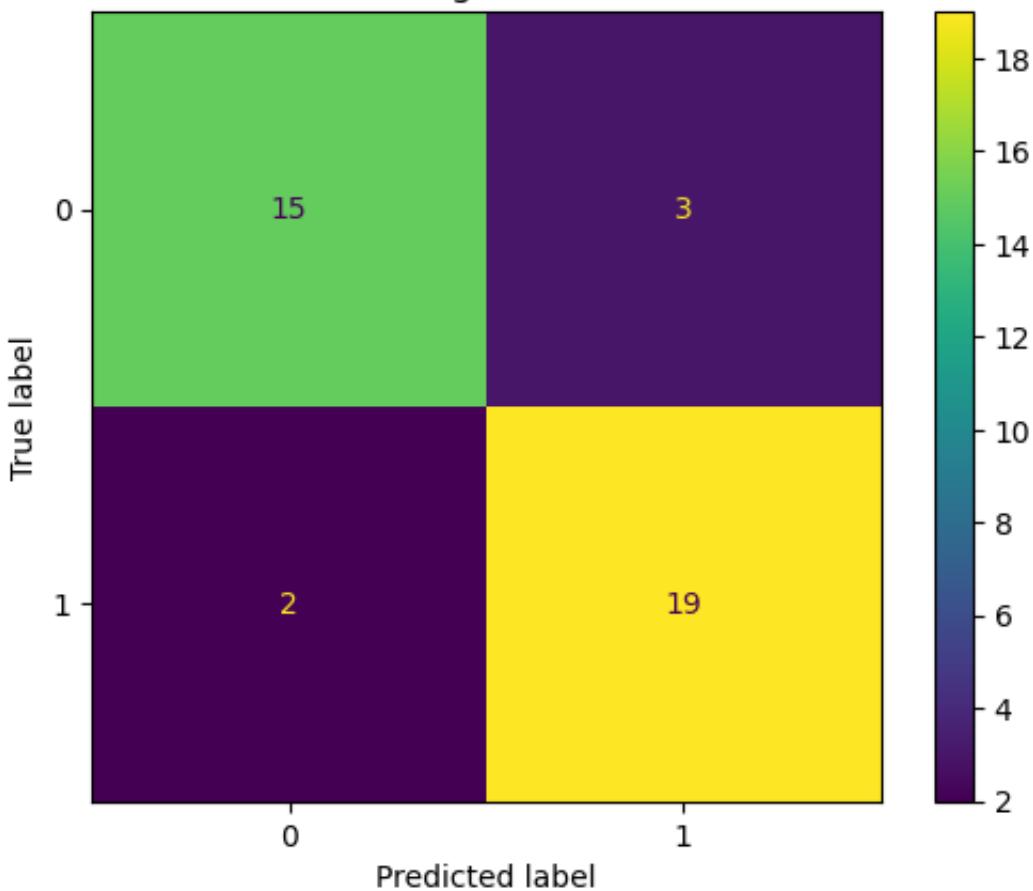
==== GradientBoostingClassifier (2D, binary) ===

Accuracy : 0.8718
 Balanced Accuracy : 0.8690
 Macro F1 : 0.8704

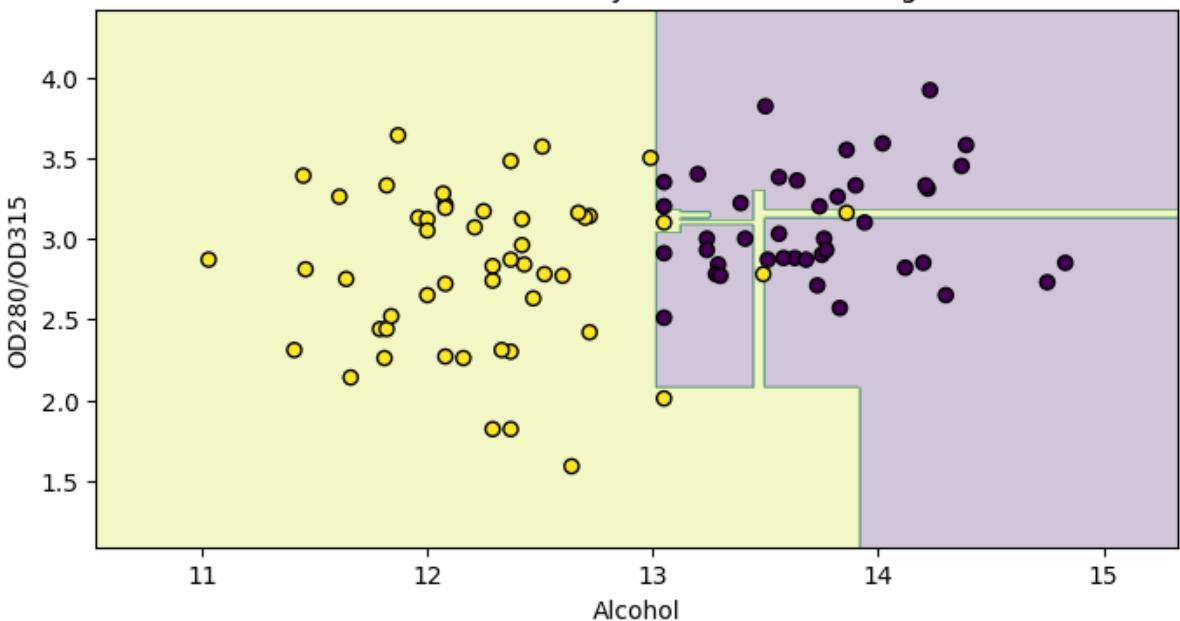
Classification Report:

	precision	recall	f1-score	support
0	0.88	0.83	0.86	18
1	0.86	0.90	0.88	21
accuracy			0.87	39
macro avg	0.87	0.87	0.87	39
weighted avg	0.87	0.87	0.87	39

GradientBoosting - Confusion Matrix



Decision Boundary - Gradient Boosting



```
In [11]: # HistGradientBoostingClassifier
# HistGradientBoosting is a more scalable implementation using histograms
# On small datasets, performance differences may be minor, but it matters

hgb = HistGradientBoostingClassifier(
    max_depth=3,
```

```
    learning_rate=0.1,
    max_iter=200,
    random_state=RANDOM_STATE
)

hgb, pred_hgb = evaluate_classifier(
    hgb, X_train, X_test, y_train, y_test,
    title="HistGradientBoostingClassifier (2D, binary)"
)

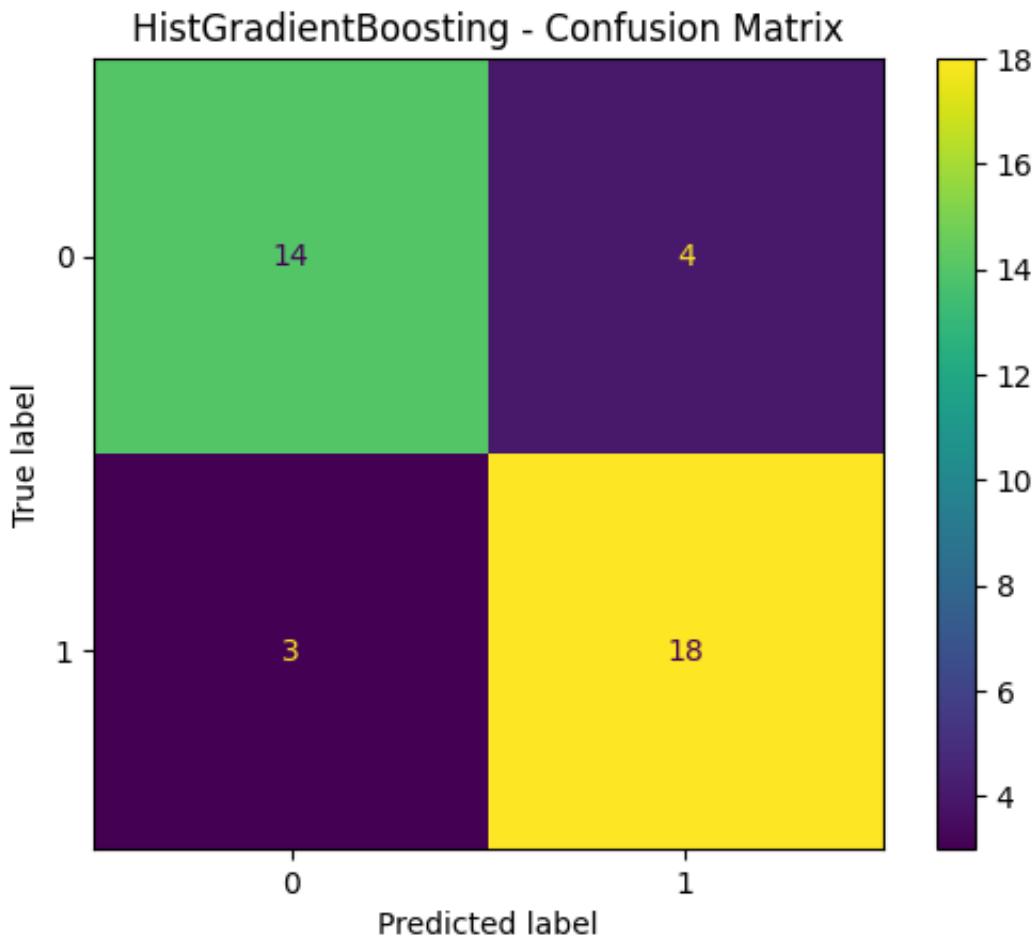
plot_confusion(y_test, pred_hgb, title="HistGradientBoosting - Confusi
```

```
== HistGradientBoostingClassifier (2D, binary) ==
```

```
Accuracy : 0.8205
Balanced Accuracy : 0.8175
Macro F1 : 0.8186
```

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.78	0.80	18
1	0.82	0.86	0.84	21
accuracy			0.82	39
macro avg	0.82	0.82	0.82	39
weighted avg	0.82	0.82	0.82	39



```
In [12]: import sys

# Install into the current Jupyter kernel's Python environment
! "{sys.executable}" -m pip install -U pip setuptools wheel
! "{sys.executable}" -m pip install xgboost
```

Requirement already satisfied: pip in ./venv/lib/python3.13/site-packages (25.3)
Requirement already satisfied: setuptools in ./venv/lib/python3.13/site-packages (80.9.0)
Requirement already satisfied: wheel in ./venv/lib/python3.13/site-packages (0.45.1)
Requirement already satisfied: xgboost in ./venv/lib/python3.13/site-packages (3.1.2)
Requirement already satisfied: numpy in ./venv/lib/python3.13/site-packages (from xgboost) (2.3.5)
Requirement already satisfied: scipy in ./venv/lib/python3.13/site-packages (from xgboost) (1.16.3)

Steps

install Homebrew

Run this in the terminal :-

```
/bin/bash -c "$(curl -fsSL  
https://raw.githubusercontent.com/Homebrew/install/HEAD/install.sh)"
```

Enter the password to your computer

==> Next steps:

- Run these commands in your terminal to add Homebrew to your PATH:

```
echo >> /Users/shivesh/.zprofile
```

```
echo 'eval "$(/opt/homebrew/bin/brew shellenv)"' >> /Users/shivesh/.zprofile
```

```
eval "$(/opt/homebrew/bin/brew shellenv)"
```

- Run brew help to get started

brew --version

brew install libomp

```
In [13]: # Now Try to Run the XGBoost Classifier  
import xgboost  
from xgboost import XGBClassifier  
xgboost.__version__
```

Out[13]: '3.1.2'

In [14]: # Part C: XGBoosting
XGBoost is an optimized, regularized gradient boosting library.
It's widely used in industry due to speed + strong performance.

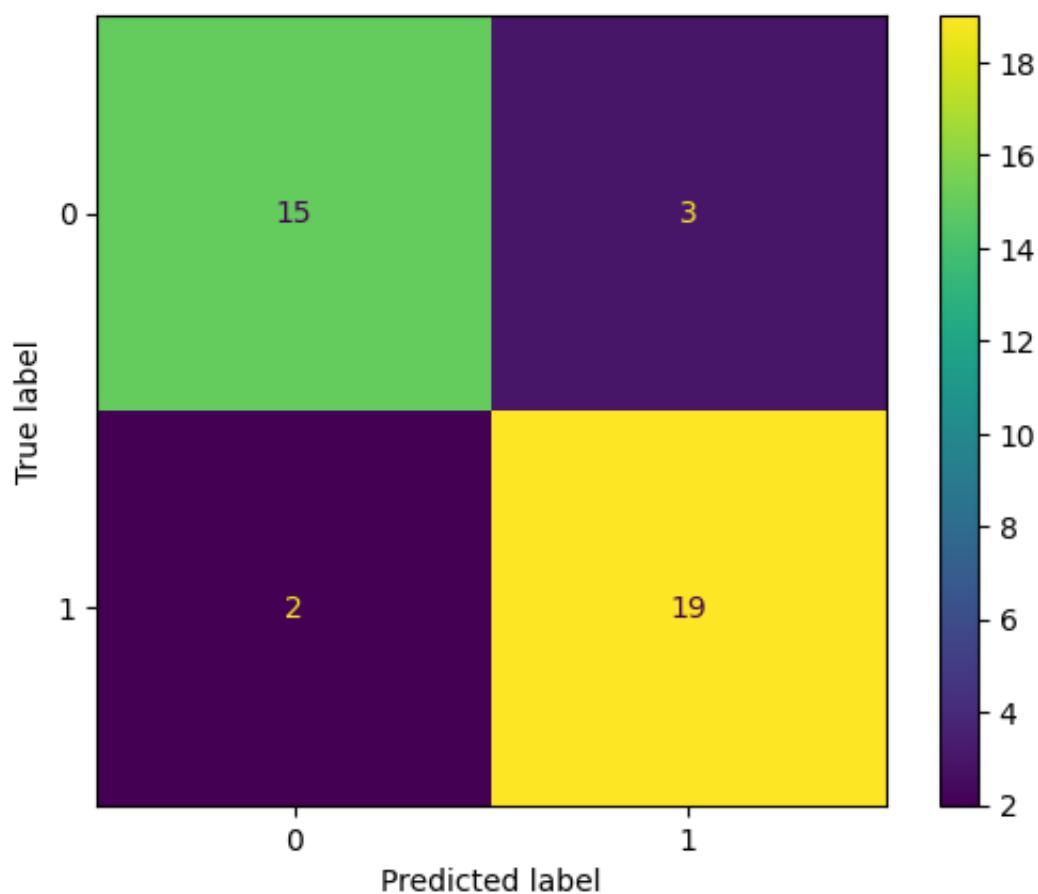
```
try:  
    from xgboost import XGBClassifier  
  
    xgb = XGBClassifier(  
        n_estimators=300,  
        learning_rate=0.05,  
        max_depth=3,  
        subsample=0.9,           # row sampling  
        colsample_bytree=0.9,    # feature sampling  
        reg_lambda=1.0,          # L2 regularization  
        reg_alpha=0.0,           # L1 regularization  
        random_state=RANDOM_STATE,  
        eval_metric="logloss"  
    )  
  
    xgb, pred_xgb = evaluate_classifier(  
        xgb, X_train, X_test, y_train, y_test,  
        title="XGBoost (2D, binary)"  
    )  
  
    plot_confusion(y_test, pred_xgb, title="XGBoost - Confusion Matrix")  
  
    plot_decision_boundary_2d(  
        xgb, X_train, y_train,  
        feature_names=("Alcohol", "OD280/OD315"),  
        title="Decision Boundary - XGBoost"  
    )  
  
except ImportError:  
    print("XGBoost is not installed.")  
    print("Install it with: !pip install xgboost")
```

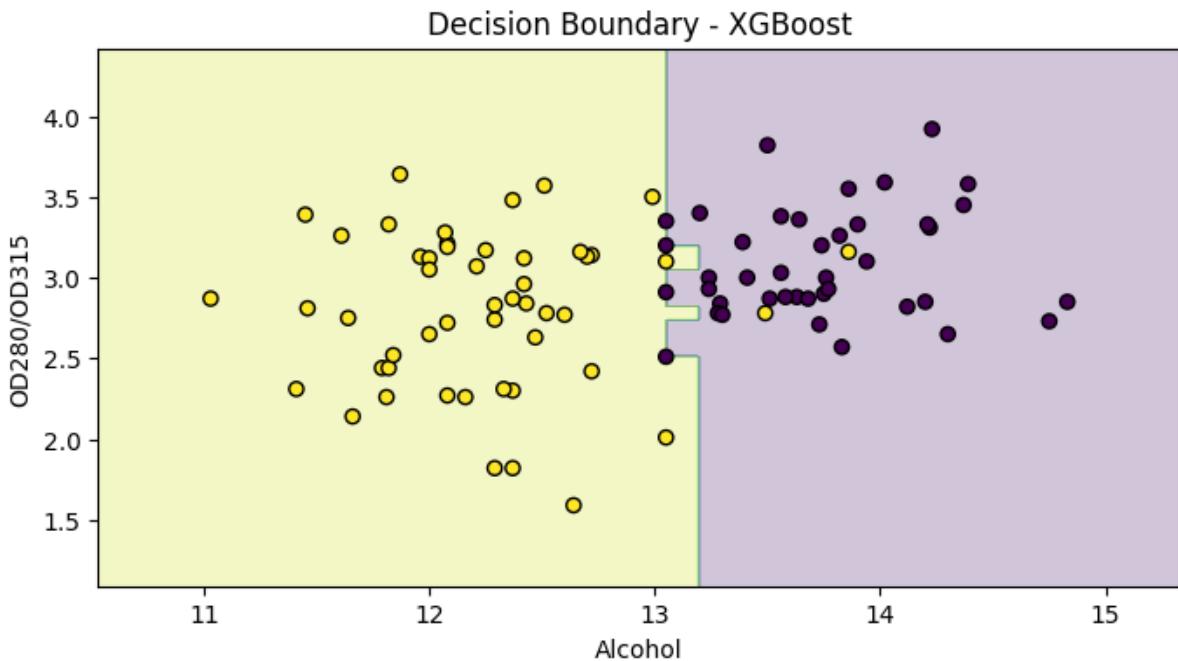
```
==== XGBoost (2D, binary) ====
Accuracy : 0.8718
Balanced Accuracy : 0.8690
Macro F1 : 0.8704
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.83	0.86	18
1	0.86	0.90	0.88	21
accuracy			0.87	39
macro avg	0.87	0.87	0.87	39
weighted avg	0.87	0.87	0.87	39

XGBoost - Confusion Matrix





```
In [15]: # Comparing All The 3 Boosting Methods
# Now we use ALL features + ALL 3 classes (multiclass classification).
X = X_full.values
y = y_full.values

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.30,
    random_state=RANDOM_STATE,
    stratify=y
)

# 1) AdaBoost for multiclass
ada_multi = AdaBoostClassifier(
    estimator=DecisionTreeClassifier(max_depth=1, random_state=RANDOM_
    n_estimators=400,
    learning_rate=0.1,
    random_state=RANDOM_STATE
)

ada_multi, pred_ada_multi = evaluate_classifier(
    ada_multi, X_train, X_test, y_train, y_test,
    title="AdaBoost (multiclass, all features)"
)

plot_confusion(y_test, pred_ada_multi, title="AdaBoost (multiclass) -"

# 2) Gradient Boosting for multiclass
gb_multi = GradientBoostingClassifier(
    n_estimators=300,
    learning_rate=0.05,
```

```
    max_depth=3,
    random_state=RANDOM_STATE
)

gb_multi, pred_gb_multi = evaluate_classifier(
    gb_multi, X_train, X_test, y_train, y_test,
    title="GradientBoostingClassifier (multiclass, all features)"
)

plot_confusion(y_test, pred_gb_multi, title="GradientBoosting (multiclass)")

# 3) XGBoost for multiclass (optional)
try:
    from xgboost import XGBClassifier

    xgb_multi = XGBClassifier(
        n_estimators=500,
        learning_rate=0.05,
        max_depth=4,
        subsample=0.9,
        colsample_bytree=0.9,
        reg_lambda=1.0,
        random_state=RANDOM_STATE,
        objective="multi:softprob",
        num_class=3,
        eval_metric="mlogloss"
    )

    xgb_multi, pred_xgb_multi = evaluate_classifier(
        xgb_multi, X_train, X_test, y_train, y_test,
        title="XGBoost (multiclass, all features)"
)

    plot_confusion(y_test, pred_xgb_multi, title="XGBoost (multiclass)")

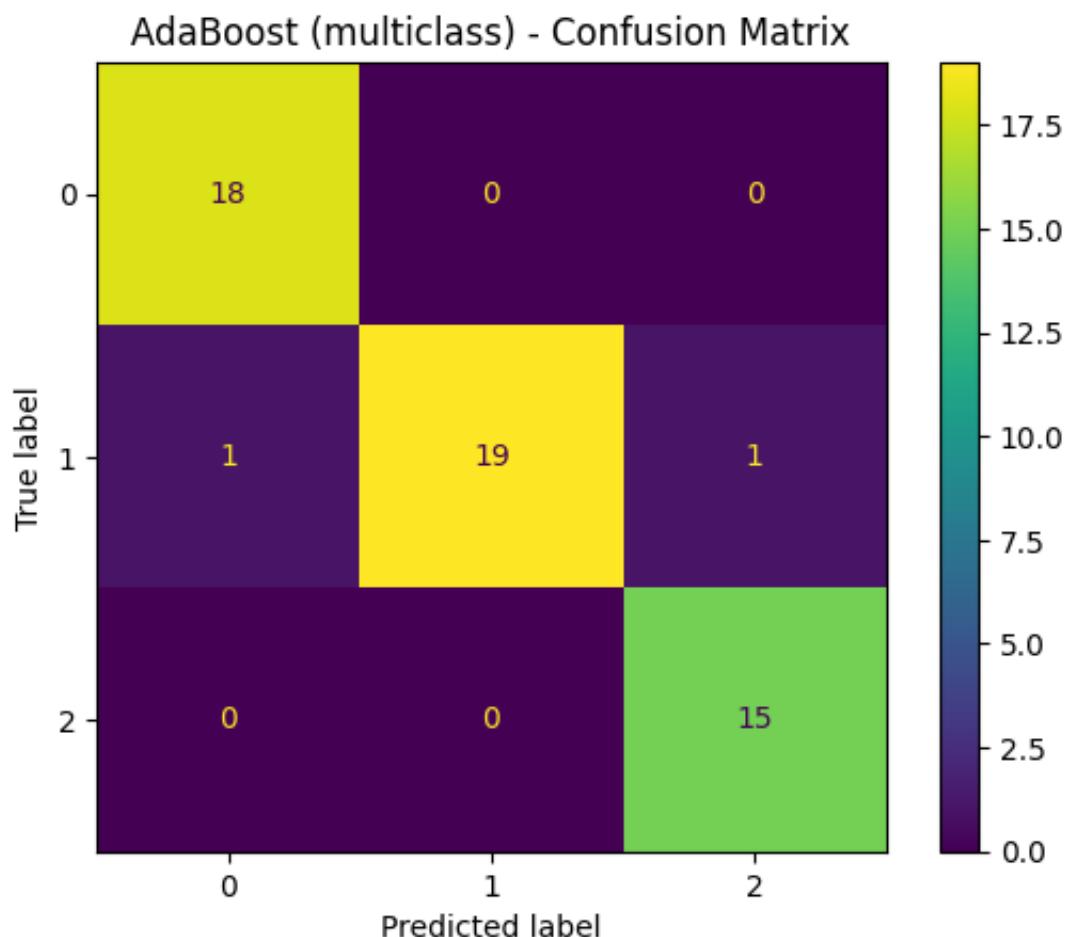
except ImportError:
    print("XGBoost not installed. Run: !pip install xgboost")
```

== AdaBoost (multiclass, all features) ==

Accuracy : 0.9630
Balanced Accuracy : 0.9683
Macro F1 : 0.9636

Classification Report:

	precision	recall	f1-score	support
0	0.95	1.00	0.97	18
1	1.00	0.90	0.95	21
2	0.94	1.00	0.97	15
accuracy			0.96	54
macro avg	0.96	0.97	0.96	54
weighted avg	0.97	0.96	0.96	54



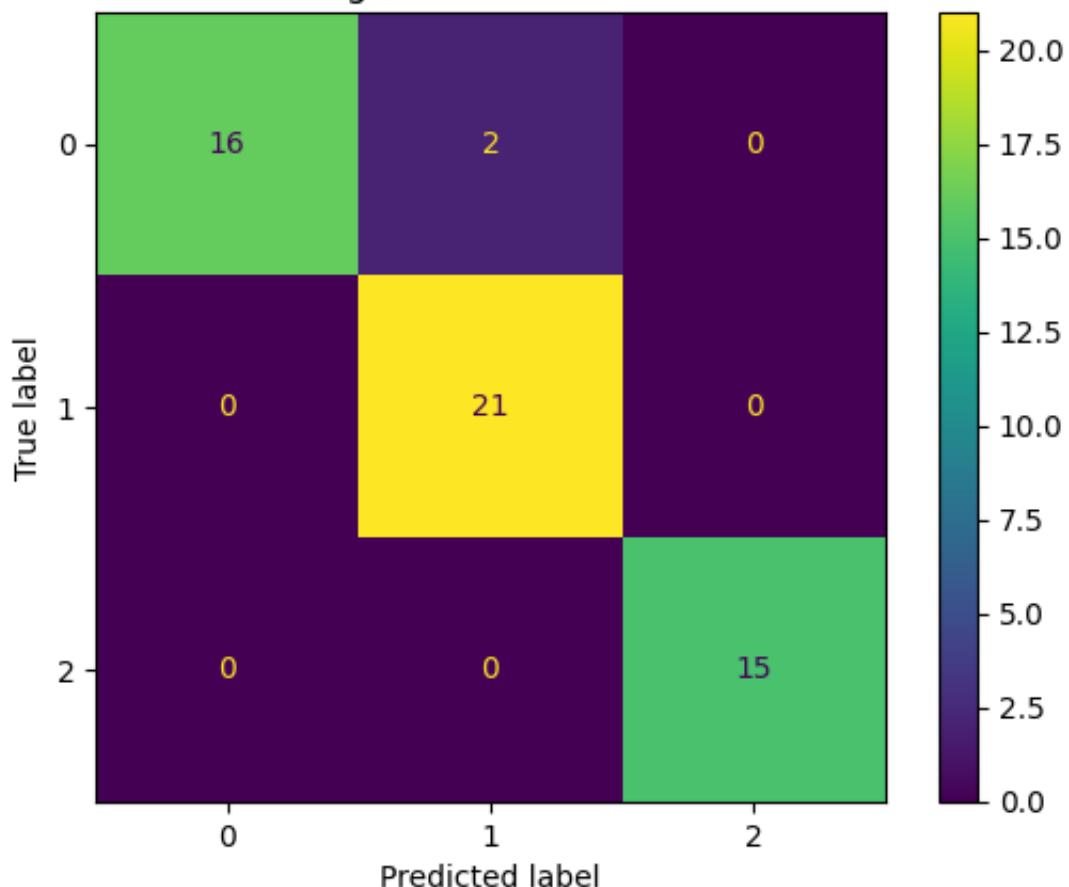
== GradientBoostingClassifier (multiclass, all features) ==

Accuracy : 0.9630
Balanced Accuracy : 0.9630
Macro F1 : 0.9652

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.89	0.94	18
1	0.91	1.00	0.95	21
2	1.00	1.00	1.00	15
accuracy			0.96	54
macro avg	0.97	0.96	0.97	54
weighted avg	0.97	0.96	0.96	54

GradientBoosting (multiclass) - Confusion Matrix

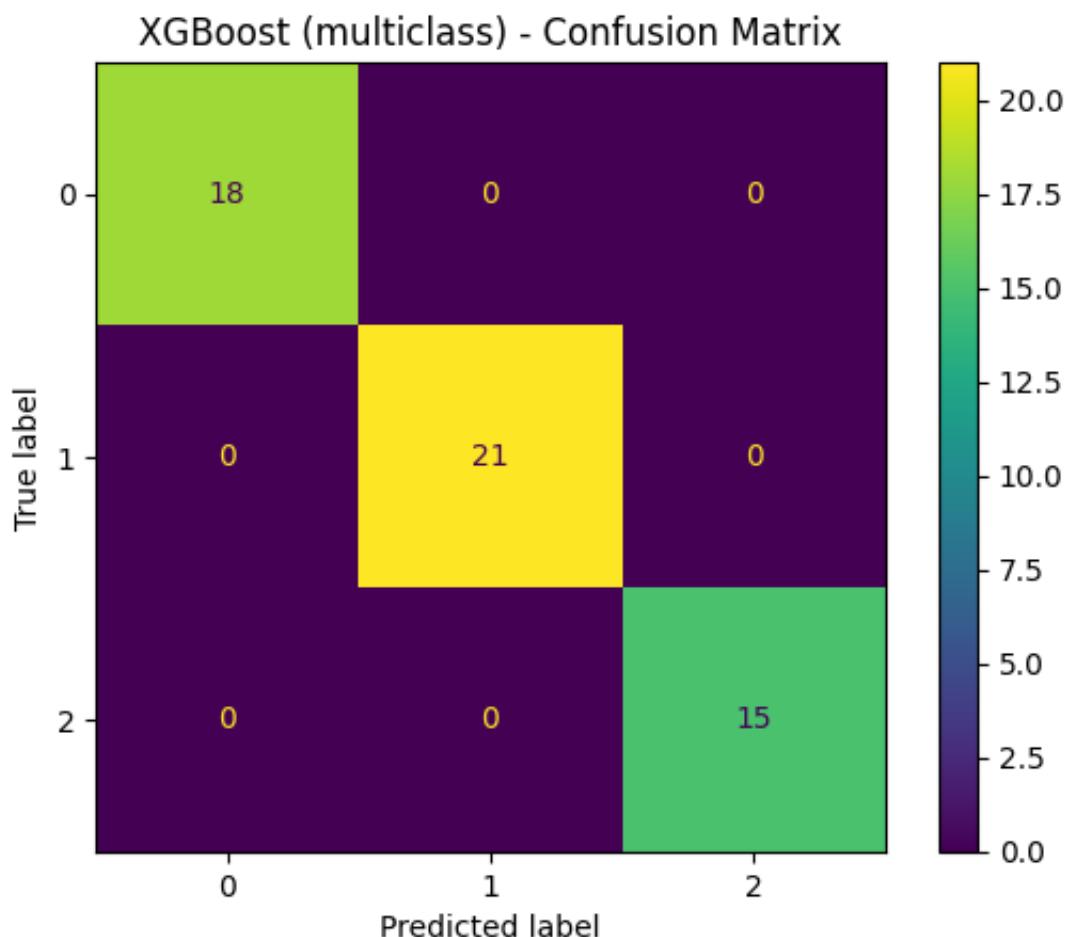


```
== XGBoost (multiclass, all features) ==
```

```
Accuracy : 1.0000
Balanced Accuracy : 1.0000
Macro F1 : 1.0000
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	18
1	1.00	1.00	1.00	21
2	1.00	1.00	1.00	15
accuracy			1.00	54
macro avg	1.00	1.00	1.00	54
weighted avg	1.00	1.00	1.00	54



Final Results & Interpretation (Boosting — AdaBoost vs Gradient Boosting vs XGBoost)

This notebook trained and evaluated **three boosting-based classifiers** on the

same dataset and splits, then compared them using:

- **Accuracy**
- **Balanced Accuracy**
- **Macro F1**
- A **Classification Report** (precision, recall, F1-score per class)
- A **Confusion Matrix** (what was predicted vs what was actually true)

Below is a complete explanation of what each output means **and how to interpret your specific results.**

1) What the Confusion Matrix is actually showing

A **confusion matrix** is a table that counts predictions:

- **Rows = True labels** (what the sample really is)
- **Columns = Predicted labels** (what our model guessed)

For a 3-class problem (classes `0`, `1`, `2`), the matrix looks like:

- Cell (0,0): number of class `0` samples correctly predicted as `0`
- Cell (0,1): number of class `0` samples wrongly predicted as `1`
- Cell (1,2): number of class `1` samples wrongly predicted as `2`
- and so on...

Diagonal cells = correct predictions **Off-diagonal cells** = mistakes (misclassifications)

So when you see a confusion matrix with most of its mass on the diagonal, that means the model is classifying well.

2) Interpreting your AdaBoost results

From your AdaBoost output (multiclass):

- **Accuracy ≈ 0.963**
- **Balanced Accuracy ≈ 0.963**
- **Macro F1 ≈ 0.964 (approx)**

What that means:

- The model predicted about **96.3% of the test samples correctly.**

- The balanced accuracy being almost the same means it is doing **fairly evenly** across classes (not just "getting the biggest class right").
- Macro F1 being close to accuracy suggests the model is not heavily failing on one class.

Confusion matrix pattern you showed:

You had **a small number of mistakes**, mainly around the middle class (class 1) being confused with others. This is common in multiclass classification when classes overlap in feature space.

Key takeaway for AdaBoost: Strong performance, but it still made a few errors—meaning it learned useful decision boundaries, but it's not perfectly separating all classes.

3) Interpreting your Gradient Boosting results

From your GradientBoosting output (multiclass):

- **Accuracy ≈ 0.963**
- **Balanced Accuracy ≈ 0.963**
- **Macro F1 ≈ 0.965 (approx)**

What that means:

This model performed **almost identically** to AdaBoost on your test set.

However, the confusion matrix suggests a *slightly different mistake pattern*:

- Typically, Gradient Boosting tends to be smoother and sometimes more stable than AdaBoost.
- You still have **a small number of off-diagonal values**, meaning a few misclassifications exist.

Key takeaway for Gradient Boosting: Comparable to AdaBoost here. Small errors remain, but overall very strong generalization.

4) Interpreting your XGBoost results (Why it shows 1.00 everywhere)

From your XGBoost output (multiclass):

- **Accuracy = 1.00**
- **Balanced Accuracy = 1.00**
- **Macro F1 = 1.00**
- Confusion matrix is perfectly diagonal (no off-diagonal values)

What that means:

On the **test set**, XGBoost predicted **every single sample correctly**.

This can happen for 3 common reasons:

(A) XGBoost is genuinely the best fit here

XGBoost is a more advanced and optimized form of gradient boosting:

- Strong regularization options (L1/L2)
- Column subsampling, row subsampling
- Better tree growth strategies
- Better numerical stability and optimization

So it can outperform scikit-learn's classic GradientBoostingClassifier.

(B) The dataset split might be "easy"

If the test subset is relatively cleanly separable, the model may achieve perfect test performance—especially with strong tree ensembles.

(C) Watch for accidental data leakage (important)

Perfect scores are awesome—but in real ML work, we always verify we didn't accidentally leak information.

Common leakage causes:

- Scaling/preprocessing done **before** train-test split
- Target variable accidentally included inside `X`
- Using `X_test` during training (even indirectly)
- Doing feature selection using the full dataset

If your workflow split first and then did preprocessing only on training data (or used a Pipeline), you are safe.

Key takeaway for XGBoost: Best result in your run (perfect test performance).

Still verify: "Is this real generalization or was there leakage?"

5) What Accuracy vs Balanced Accuracy vs Macro F1 mean (simple but exact)

Accuracy

[Accuracy = $\frac{\text{Total correct predictions}}{\text{Total predictions}}$]

- Good overall metric
- But can be misleading if one class dominates

Balanced Accuracy

Balanced accuracy is the average recall across classes: [BalancedAccuracy = $\frac{\text{Recall}_0 + \text{Recall}_1 + \text{Recall}_2}{3}$]

- Helps when classes are imbalanced
- Answers: "Am I treating all classes fairly?"

Macro F1

Macro F1 averages F1 across classes equally: [MacroF1 = $\frac{\text{F1}_0 + \text{F1}_1 + \text{F1}_2}{3}$]

- Penalizes models that fail on a minority/rare class
- A strong "fair" overall metric for multiclass

Because your Accuracy, Balanced Accuracy, and Macro F1 are all close, your models are likely doing well **across all classes**, not just one.

6) Precision, Recall, F1 Score

For each class:

Precision

"When the model predicts class X, how often is it correct?"

High precision = few false positives.

Recall

"Out of all true class X samples, how many did the model catch?"

High recall = few false negatives.

F1-score

F1 balances precision and recall: [$F1 = \frac{2 \cdot \text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$]

So:

- If precision is high but recall low → model is too conservative
 - If recall high but precision low → model predicts that class too easily
 - High F1 means both are strong.
-

7) Practical summary

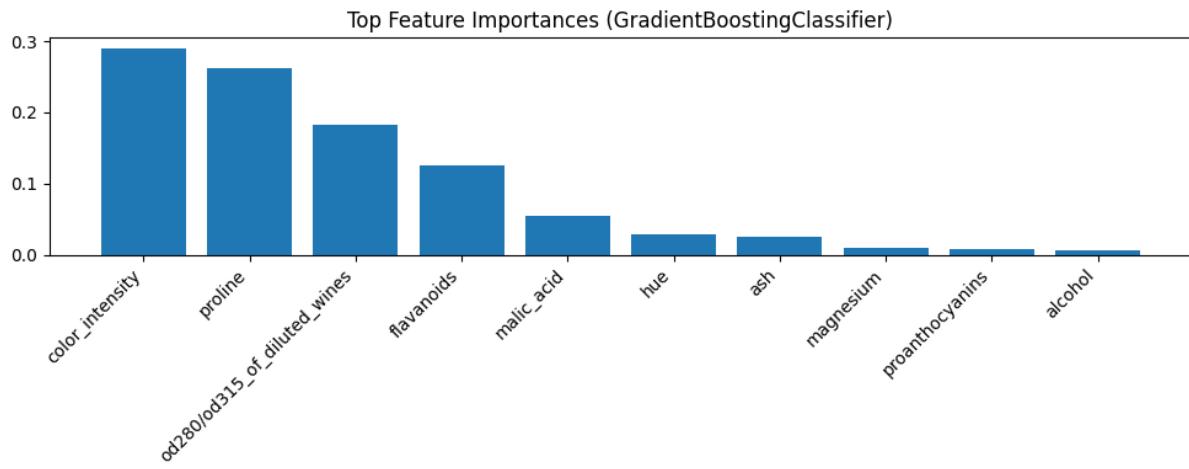
- **AdaBoost** performed very strongly (~96%), with a few misclassifications.
 - **Gradient Boosting** matched AdaBoost performance (~96%) and showed similar generalization.
 - **XGBoost** achieved **perfect classification** on the test set in this run, likely due to:
 - stronger regularization + robust boosting implementation,
 - and/or a cleanly separable test split.
 - Confusion matrices confirm:
 - AdaBoost/GB had small off-diagonal errors,
 - XGBoost was fully diagonal.
-

```
In [16]: # Feature Importance for Gradient Boosting (Scikit-Learn)
# Many tree-based models expose "feature_importances_".
# This gives a rough idea of which features are most useful.
# (Not the same as causal importance.)

importances = gb_multi.feature_importances_
indices = np.argsort(importances)[-1:] # sort descending

top_k = 10
top_indices = indices[:top_k]

plt.figure(figsize=(10, 4))
plt.bar(range(top_k), importances[top_indices])
plt.xticks(range(top_k), np.array(feature_names)[top_indices], rotation=90)
plt.title("Top Feature Importances (GradientBoostingClassifier)")
plt.tight_layout()
plt.show()
```

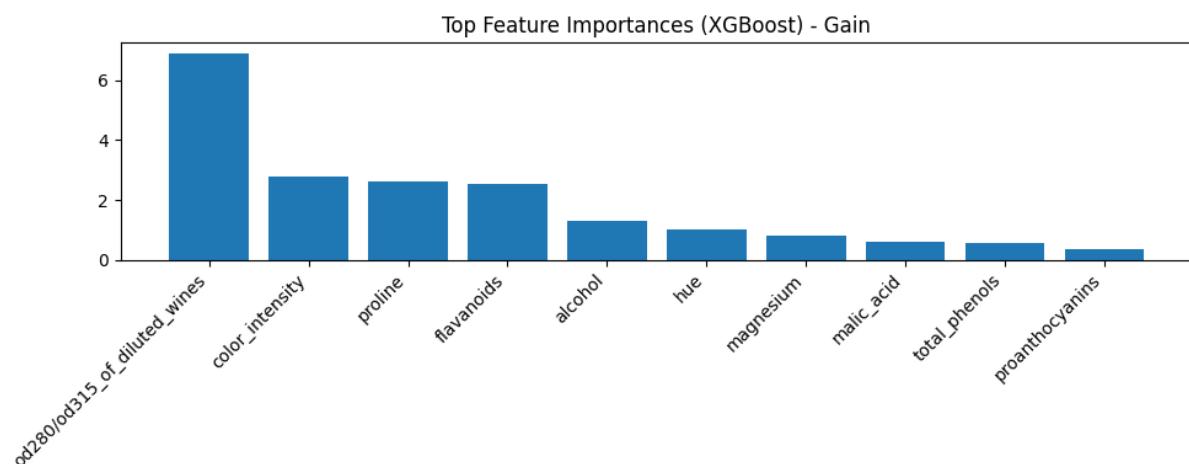


```
In [17]: # Feature Importance For XGBoosting
try:
    booster = xgb_multi.get_booster()
    score = booster.get_score(importance_type="gain") # "gain" is a convention

    # Convert dict -> DataFrame for sorting
    df_imp = pd.DataFrame(score.items(), columns=["feature", "gain"])
    df_imp["feature_index"] = df_imp["feature"].str.replace("f", "", regex=True)
    df_imp["feature_name"] = df_imp["feature_index"].apply(lambda i: f"od280/od315_of_diluted_wines" if i == 0 else i)
    df_imp = df_imp.sort_values("gain", ascending=False).head(10)

    plt.figure(figsize=(10, 4))
    plt.bar(df_imp["feature_name"], df_imp["gain"])
    plt.xticks(rotation=45, ha="right")
    plt.title("Top Feature Importances (XGBoost) - Gain")
    plt.tight_layout()
    plt.show()

except Exception as e:
    print("XGBoost importance plot skipped (either not installed or missing dependencies)")
    print("Reason:", e)
```



In []:

